

NPV-Optimal SMB Underwriting under Selection Bias, with Calibrated Risk and Interventional Counterfactuals

Team Vector Ventures

June 2026

1 Problem framing & assumptions violated

Every assumption a credit scorecard leans on—random labels, population overlap, a stationary window, honest inputs—breaks on this book; we name each break and respond to it rather than modelling straight through it.

We are re-underwriting a historical SMB loan book to maximize realized portfolio net present value (NPV)—total dollars earned, not classification accuracy.

Labels are missing-not-at-random behind a sharp gate. Outcomes exist only for loans the prior lender approved, and approval is, to numerical precision, a hard cutoff `prior_underwriter_score` ≥ 0.273 : a propensity model separates approved from declined with AUC 1.00 and *zero* score overlap. So the label is selected on a deterministic function of an observed covariate—textbook sample-selection bias. Inverse-propensity reweighting is infeasible here by a positivity violation, *not by our choice*: a propensity model on the ordinary covariates alone still hits AUC 1.00, so approved and declined applicants share no common feature support. Only local extrapolation near the cutoff is available, and we widen intervals on declines rather than pretend we can reweight.

The split is forward in time. Training spans an 18-month book ending the day before the 13-week scoring window opens, over which the default rate drifts 17.5% $\rightarrow$ 20.6%. We learn default *timing* on the history but recalibrate the *level* in-window—the history tells us the shape, only the window tells us the rate.

Missingness is information, not absence. “Never declined elsewhere” (null) loans default at 14.3% versus 21.3% when present; no-bank-feed loans at 19.1% versus 16.7%—a blank field predicts the outcome, so we encode it as a first-class indicator rather than impute it away. **Self-reports are optimistically biased:** applicants who overstate revenue by more than $1.5\times$ default at 38.6% versus 16.1%, so the *gap* between stated and bank-feed revenue—not the stated figure—is the signal, central to our causal treatment of `stated_annual_revenue` below.

2 Methodology

One discrete-time model serves all three deliverables: a calibrated probability that a loan defaults, times a per-band curve for when it defaults, read out as an NPV sign for the approve/decline call.

One shared discrete-time model. The brief’s NPV depends on the *day* of default (a day-5 default loses roughly the principal; a day-55 default nearly breaks even), and Deliverable B is the cumulative-default curve. Both consume one timing model. We write each loan’s cumulative-default curve as $F_i(t) = \text{PD}_i \cdot S_{b(i)}(t)$, where $\text{PD}_i = \Pr(y_i=1 \mid \mathbf{x}_i)$ is a calibrated probability of default by day 90 and $S_b(t)$ is a normalized timing shape ($S_b(90)=1$) learned per credit band b . In words: *how likely* a default is, kept separate from *when* it arrives.

PD model. A five-fold `GroupKFold` (by `business_id`) `LightGBM` ensemble over 55 engineered features (out-of-fold AUC 0.774, Brier 0.117), isotonic-calibrated. Features are deliberately buffer-centric: because loans are $\sim 1.3\%$ of annual revenue, default is cash-buffer bound rather than revenue bound, so debt-service coverage, buffer-to-payment, overdrafts and credit utilization dominate.




Figure 1: Pooled cumulative default rate by loan-age week: a concave rise through week 9, a flat stretch (weeks 9–12), and the day-90 open-balance cliff at week 13. Deliverable B reproduces this shape per cohort; a single smooth survival curve cannot.

Two-mode hazard for timing (Figure 1). 77.5% of defaults are missed-draw events over days 3–60; *zero* occur on days 61–89; the remaining 22.5% spike at day 90 (open-balance sweep). We model the two modes as separate heads: PD above, plus $\text{d90_frac}_i = \Pr(\text{day-90 sweep} \mid \text{default}, \mathbf{x}_i)$ (out-of-fold AUC 0.72). Per loan, $\text{CDR}_i(a) = \text{PD}_i(1 - \text{d90_frac}_i)F_{\text{early}}^{b(i)}(a)$ for $a < 13$ and $\text{CDR}_i(13) = \text{PD}_i$, where F_{early}^b is the band-conditional cumulative-default fraction within days 3–60. A’s $\mathbb{E}[\text{NPV}]$ uses $\mathbb{E}[t^* \mid \text{default}, i] = (1 - \text{d90_frac}_i)\mathbb{E}[t \mid \text{early}, b(i)] + \text{d90_frac}_i \cdot 90$ (exact under linear-in- t NPV).

Deliverable A: decide on the NPV sign. Approve iff $\mathbb{E}[\text{NPV}_i] > 0$ under brief NPV: $F_i + R_i r(T/365)$ if repaid; $F_i + D_i(t_i^* - 1) + \text{rec}_i - R_i$ if default at day t_i^* (fee + pre- t^* draws + recovery – principal). Break-even PD ≈ 0.51 under two-mode timing. We shift the decision PD by $\kappa \hat{\sigma}_i$ ($\hat{\sigma}_i$ fold disagreement, $\kappa=2.25$ by cross-fit on realized val P&L: OOF \$3.89M vs \$3.91M at $\kappa=0$), declining the most uncertain near-break-even loans. We approve 67% of the 13,306 scored applicants; on labeled val the book returns $1.83\times$ the prior underwriter (capture 0.60 of oracle).

Deliverable B: shape $\times$ level. $\text{CDR}_{w,a}$ averages the per-loan $\text{CDR}_i(a)$ above, monotone by construction with the day-90 cliff at $a=13$. The PD model has no cohort signal (training predates the window), so we shrink each cohort’s *level* toward the realized val rate (pseudo-count 15) and *shape* toward the model band shape (Dirichlet $c=50$); we also rescale per-cohort PD for the two biased cohorts (5, 13; $|\text{bias}| > 0.04$). Val and test share the same 13 calendar weeks, so the empirical shape transfers. MAE 0.013 $\rightarrow$ 0.003.

3 Causal reasoning & counterfactual methodology

The judges’ counterfactual queries ask what an intervention causes, not what a feature predicts—so we classify every feature by whether editing it changes reality, and scale the effect accordingly.

The counterfactual target is $p_q^{\text{cf}} = \Pr(y_{i_q}=1 \mid \text{do}(f_q=v_q), \mathbf{X}_{i_q, -f_q}=\mathbf{x}_{i_q, -f_q})$ versus the observational $\Pr(y \mid f=v, \mathbf{X}_{-f})$: the first asks what happens if we *set* f to v , the second describes applicants who happen to *have* $f=v$, and they diverge exactly when f is confounded. Our **feature registry** keeps the edit honest: each engineered feature declares its parents, so setting a raw feature recomputes exactly its downstream features and nothing else—never an internally-contradictory applicant carrying a \$20k loan with a \$10k loan’s payments. For example `do(requested.amount)` recomputes the daily-payment, buffer-to-payment, leverage and ask-to-revenue features; the provided `requested.amount_to_observed_revenue` is never reused (no stale-descendant bug).

We classify the intervenable features by causal status. *Mechanical causes* (`requested.amount`): changing the number changes the borrower’s actual payment burden, so we perturb and propagate it through the registry at full strength. *Confounded proxies*—bureau and bank-feed symptoms of latent health (utilization, delinquency, revenue and cash signals)—are symptoms as much as causes: editing the thermometer does not cure the fever, so we keep only the estimated causal fraction $\hat{\lambda} = \beta_{\text{adj}}/\beta_{\text{naive}}$ (sibling-adjusted logistic; per-family $\approx 0.52/0.51/0.64$ for bureau/bank-feed/behavioral), averaged with the heuristic, falling back where $\hat{\lambda} \notin (0, 1)$. Conditioning sets are derived *algorithmically* from an explicit




Figure 2: SHAP summary for the PD model: mean $|\text{SHAP}|$ ranking with per-feature direction (red high / blue low). High delinquency, utilization and volatility raise PD; high debt-service coverage and cash buffer lower it.




Figure 3: The explicit DAG. Latent health H confounds every proxy (faint), so we adjust for a treatment’s co-proxies of H ; the three defended mechanistic edges (red) make invoice delinquency/overdrafts descendants of cash and payroll a descendant of revenue, which graph reachability excludes from those treatments’ adjustment sets. The FDR partial-correlation graph ($q=0.05$, $-0.78/+0.50$ edges) and PC/GES corroborate this structure.

DAG by the backdoor criterion—condition on a treatment’s co-proxies of latent health, exclude its graph *descendants* (Figure 3). The FDR partial-correlation graph flags “siblings” that are really descendants—invoice delinquency and overdrafts follow low cash (-0.78), payroll regularity follows revenue ($+0.50$); PC and GES corroborate. Since conditioning on a descendant blocks the true causal path, excluding it by graph reachability corrects $\hat{\lambda}_{\text{cash}}$ from an implausible 0.03 to 0.44. *Self-reports*—**stated_annual_revenue**, **stated_time_in_business**—are returned at their observational PD: a different number on the form changes nothing about the business. Overstating revenue *predicts* default (38.6% vs 16.1%), but *forcing* a different statement causes nothing, so the interventional effect is ≈ 0 —formally $\Pr(y \mid \text{do}(\text{stated}), \mathbf{X}_{\text{stated}}) = \Pr(y \mid \mathbf{X}_{\text{stated}})$, the causally correct answer, not naive re-prediction. *Immutable/proxy-by-fiat* features (**sector**, **geography_region**, **vintage_years**, **has_linked_bank_feed**) are not levers but are queried anyway; we perturb the model, shrink, widen, and flag them.

Regulator defense. Drivers are explained with SHAP over an inspectable, de-correlated feature set (Figure 2). The leading drivers—invoice-payment delinquency, credit utilization, revenue volatility, requested amount, overdrafts, and the affordability ratios—are all recognized credit/affordability signals in the expected direction, none a prohibited or proxy-discriminatory attribute. Several queried values lie outside observed support (e.g. **prior_loans_count** is in-support only 29% of the time); these receive wider intervals and an explicit extrapolation caveat.

Model-class check. A transparent logistic regression on the same engineered features matches the

gradient-boosted models (out-of-fold AUC 0.775 vs 0.770, equal Brier) and is in fact *more* robust to the train-to-window drift (validation AUC 0.759 vs 0.751)—evidence the engineered space is near-linearly sufficient and the drivers directly interpretable. Its coefficients agree in sign with SHAP on 8/10 top features (the two flips, `buffer_to_payment` and `debt_to_revenue`, are collinearity-driven—why we read SHAP, not raw coefficients). We rejected sequence/transformer models: no sequential structure, low-cardinality categoricals, ≈ 0.77 -AUC data-limited discrimination, and calibration favours the simpler stack.

4 Calibration & uncertainty quantification

We fix the probability level on the only data drawn from the scoring window, then size every interval by measured out-of-sample error rather than in-sample optimism.

Point PD. Isotonic calibration is *refit on in-window validation* to correct the 17.5%→20.6% train-to-window drift (power-prior $a_0 \rightarrow 0$ limit); it stays valid on the approved manifold and is extrapolated on declines.

Intervals (A, C). Additive, $p \pm \alpha z \sigma$ (σ = fold disagreement; $\alpha=1.1$ by cross-fit on val, target 0.90). On labeled val: decile coverage 1.00 / per-cohort 12/13 at mean width 0.064—the cross-fit reading hits 0.90 exactly, so the shipped over-coverage is the post-isotonic-refit tail margin, not slack.

Trajectory intervals (B). Within-cohort bootstrap combined with a conformal band against the true val cohort trajectories; coverage ≈ 0.94 at width 0.048. Out-of-support counterfactuals (C) are widened further (Table 1).

Deliverable	Metric	Value
A — PD model	out-of-fold AUC / Brier	0.774 / 0.117
A — d90 head	out-of-fold AUC / val base rate	0.72 / 0.21
A — policy	approve rate (scored) / P&L vs. prior underwriter	0.67 / 1.83×
A — intervals	decile / per-cohort coverage / mean width	10/10 / 12/13 / 0.064
B — trajectory	MAE vs. realized val CDR / coverage / width	0.003 / 0.94 / 0.048

Table 1: Validation evidence (in-window). P&L and B error use realized validation outcomes; interval coverage is out-of-sample (cross-fit / conformal).

5 Limitations & what we’d do differently

Our honest gaps trace to one fact—we never observe outcomes for declined applicants—so the decline and counterfactual estimates rest on extrapolation we cannot check.

The unverifiable core. All C effects lack labels (the queried applicants were never decided), and declines never overlap the approved manifold (§1)—so reject inference is infeasible, only RD-style local extrapolation remains (deferred), and decline-side PD, part of A’s score, stays unchecked. Self-report and immutable counterfactuals stay heuristic; confounded proxies use the DAG-derived $\hat{\lambda}$ above, but latent H breaks causal sufficiency—PC/GES corroborate the structure, they do not prove it.

Our val-tuned knobs ($\kappa=2.25$, level pseudo-count 15, shape $c=50$, per-cohort PD scaling for cohorts 5 and 13) transfer to test only as far as the windows match (same 13 weeks; per-cohort $n \approx 150$ adds variance). The κ rule absorbs the flagged over-approval, but a residual oracle gap remains (≈ 0.60 of perfect-foresight P&L), irreducible without label foresight on declines.

With another day: sharp-RD local-linear PD anchor across the 0.273 cutoff to calibrate decline-side PD (identifiable via continuity even though positivity fails), cohort_week as a feature in the d90 head (would tighten the per-cohort d90-share residual, ~ 7 pp on cohort 13), and a sealed val-tuning holdout to honestly trade off the seven val-fit calibrations against single-split overfit risk.